

# Bios 6301: Assignment 5

*Elizabeth Sigworth*

*Tuesday 15 November*

*Due Tuesday, 15 November, 1:00 PM*

$5^{n=\text{day}}$  points taken off for each day late.

50 points total.

Submit a single knitr file (named `homework5.rmd`), along with a valid PDF output file. Inside the file, clearly indicate which parts of your responses go with which problems (you may use the original homework document as a template). Add your name as `author` to the file's metadata section. Raw R code/output or word processor files are not acceptable.

Failure to name file `homework5.rmd` or include author name may result in 5 points taken off.

## Question 1

### 24 points

Import the HAART dataset (`haart.csv`) from the GitHub repository into R, and perform the following manipulations: (4 points each)

```
url1 <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart <- read.csv(url1, stringsAsFactors=FALSE)
```

1. Convert date columns into a usable (for analysis) format. Use the `table` command to display the counts of the year from `init.date`.

```
haart$init.date <- as.Date(haart$init.date, format="%m/%d/%y")
haart$last.visit <- as.Date(haart$last.visit, format="%m/%d/%y")
haart$date.death <- as.Date(haart$date.death, format="%m/%d/%y")
table(format(haart$init.date, '%Y'))
```

```
##
## 1998 2000 2001 2002 2003 2004 2005 2006 2007
##    1    5   17   60  270  292  207  104   44
```

2. Create an indicator variable (one which takes the values 0 or 1 only) to represent death within 1 year of the initial visit. How many observations died in year 1?

```
for (i in 1:nrow(haart)) {
  haart$one.year[i] <- ifelse(abs(unclass(difftime(haart$init.date[i],
                                                    haart$date.death[i], units='days'))[1]) > 365, 0, 1)
}
year.one <- sum(haart$one.year, na.rm = TRUE)
```

In this data, 92 observations died in year 1.

- Use the `init.date`, `last.visit` and `death.date` columns to calculate a followup time (in days), which is the difference between the first and either the last visit or a death event (whichever comes first). If these times are longer than 1 year, censor them (this means if the value is above 365, set followup to 365). Print the quantile for this new variable.

```
for (i in 1:nrow(haart)){
  if(is.na(haart$date.death[i]) == TRUE ) {
    difference <- unclass(difftime(haart$last.visit[i], haart$init.date[i], 'days'))[1]
    haart$follow.up[i] <- min(365,difference)
  }
  else if(is.na(haart$date.death[i]) == FALSE && is.na(haart$last.visit[i])==TRUE){
    difference <- unclass(difftime(haart$date.death[i], haart$init.date[i], 'days'))[1]
    haart$follow.up[i] <- min(365,difference)
  }
  else {
    first <- min(haart$last.visit[i],haart$date.death[i])
    difference <- unclass(difftime(first, haart$init.date[i], 'days'))[1]
    haart$follow.up[i] <- min(365,difference)
  }
}
quantile(haart$follow.up)
```

```
##      0%      25%      50%      75%     100%
##    0.00 320.75 365.00 365.00 365.00
```

- Create another indicator variable representing loss to followup; this means the observation is not known to be dead but does not have any followup visits after the first year. How many records are lost-to-followup?

```
for (i in 1:nrow(haart)){
  if(is.na(haart$date.death[i]) && haart$follow.up[i] < 365){
    haart$lost[i] <- 1
  }
  else {
    haart$lost[i] <- 0
  }
}
lost.to <- sum(haart$lost, na.rm=TRUE)
```

There were 173 records lost-to-followup.

- Recall our work in class, which separated the `init.reg` field into a set of indicator variables, one for each unique drug. Create these fields and append them to the database as new columns. Which drug regimen are found over 100 times?

```
init.reg <- as.character(haart[, 'init.reg'])
haart[['init.reg_list']] <- strsplit(init.reg, ",")
(all_drugs <- unique(unlist(haart$init.reg_list)))
```

```
## [1] "3TC" "AZT" "EFV" "NVP" "D4T" "ABC" "DDI" "IDV" "LPV" "RTV" "SQV"
## [12] "FTC" "TDF" "DDC" "NFV" "T20" "ATV" "FPV"
```

```
(unique_drugs <- unique(unlist(haart$init.reg_list)))
```

```
## [1] "3TC" "AZT" "EFV" "NVP" "D4T" "ABC" "DDI" "IDV" "LPV" "RTV" "SQV"
## [12] "FTC" "TDF" "DDC" "NFV" "T20" "ATV" "FPV"
```

```
reg_drugs <- matrix(FALSE, nrow=nrow(haart), ncol=length(all_drugs))
for(i in seq_along(all_drugs)) {
  reg_drugs[,i] <- sapply(haart$init.reg_list, function(x) all_drugs[i] %in% x)
}
reg_drugs <- data.frame(reg_drugs)
names(reg_drugs) <- all_drugs
haart_merged <- cbind(haart, reg_drugs)
for (i in 17:34){
  for (j in 1:nrow(haart_merged)){
    if(haart_merged[j,i]==TRUE){
      haart_merged[j,i] <- colnames(haart_merged)[i]
    }
    else {
      haart_merged[j,i] <- NA
    }
  }
}

apply(X = haart_merged[,17:34],2,table)
```

```
## 3TC AZT EFV NVP D4T ABC DDI IDV LPV RTV SQV FTC TDF DDC NFV T20 ATV FPV
## 973 794 516 358 146 56 38 27 31 79 29 8 10 1 8 1 2 2
```

The individual drugs 3TC, AZT, EFV, NVP, and D4T are each found over 100 times.

6. The dataset `haart2.csv` contains a few additional observations for the same study. Import these and append them to your master dataset (if you were smart about how you coded the previous steps, cleaning the additional observations should be easy!). Show the first five records and the last five records of the complete (and clean) data set.

```
url2 <- "https://raw.githubusercontent.com/fonnesbeck/Bios6301/master/datasets/haart2.csv"
haart2 <- read.csv(url2, stringsAsFactors=FALSE)
haart2$init.date <- as.Date(haart2$init.date, format="%m/%d/%y")
haart2$last.visit <- as.Date(haart2$last.visit, format="%m/%d/%y")
haart2$date.death <- as.Date(haart2$date.death, format="%m/%d/%y")
for (i in 1:nrow(haart2)) {
  haart2$one.year[i] <- ifelse(abs(unclass(difftime(haart2$init.date[i], haart2$date.death[i],
                                                    units='days'))[1]) > 365, 0, 1)
}
for (i in 1:nrow(haart2)){
  if(is.na(haart2$date.death[i]) == TRUE) {
    difference <- unclass(difftime(haart2$last.visit[i], haart2$init.date[i], 'days'))[1]
    haart2$follow.up[i] <- min(365,difference)
  }
  else {
    difference <- unclass(difftime(haart2$date.death[i], haart2$init.date[i], 'days'))[1]
```

```

    haart2$follow.up[i] <- min(365,difference)
  }
}
for (i in 1:nrow(haart2)){
  if(is.na(haart2$date.death[i]) && unclass(difftime(haart2$last.visit[i],
                                                    haart2$init.date[i], 'days'))[1] < 365){

    haart2$lost[i] <- 1
  }
  else {
    haart2$lost[i] <- 0
  }
}

init.reg <- as.character(haart2[, 'init.reg'])
haart2[['init.reg_list']] <- strsplit(init.reg, ",")

reg_drugs <- matrix(FALSE, nrow=nrow(haart2), ncol=length(all_drugs))
for(i in seq_along(all_drugs)) {
  reg_drugs[,i] <- sapply(haart2$init.reg_list, function(x) all_drugs[i] %in% x)
}
reg_drugs <- data.frame(reg_drugs)
names(reg_drugs) <- all_drugs
haart2_merged <- cbind(haart2, reg_drugs)

for (i in 17:34){
  for (j in 1:nrow(haart2_merged)){
    if(haart2_merged[j,i]==TRUE){
      haart2_merged[j,i] <- colnames(haart2_merged)[i]
    }
    else {
      haart2_merged[j,i] <- NA
    }
  }
}

haart_final <- rbind(haart_merged,haart2_merged)
haart_final[c(1:5,1000:1004),]

```

```

##      male      age aids cd4baseline      logvl  weight hemoglobin
## 1      1 25.00000    0      NA      NA      NA      NA
## 2      1 49.00000    0     143      NA 58.0608      11
## 3      1 42.00000    1     102      NA 48.0816       1
## 4      0 33.00000    0     107      NA 46.0000      NA
## 5      1 27.00000    0      52 4.000000      NA      NA
## 1000    0 40.00000    1     131      NA 46.2672       8
## 1001    0 27.00000    0     232      NA      NA      NA
## 1002    1 38.72142    0     170      NA 84.0000      NA
## 1003    1 23.00000   NA     154 3.995635 65.5000      14
## 1004    0 31.00000    0     236      NA 45.8136      NA
##      init.reg  init.date last.visit death date.death one.year follow.up
## 1      3TC,AZT,EFV 2003-07-01 2007-02-26    0      <NA>      NA      365
## 2      3TC,AZT,EFV 2004-11-23 2008-02-22    0      <NA>      NA      365

```

```
## 3      3TC,AZT,EFV 2003-04-30 2005-11-21      1 2006-01-11      0      365
## 4      3TC,AZT,NVP 2006-03-25 2006-05-05      1 2006-05-07      1      41
## 5      3TC,D4T,EFV 2004-09-01 2007-11-13      0      <NA>      NA      365
## 1000   3TC,D4T,NVP 2003-07-03 2008-02-29      0      <NA>      NA      365
## 1001   3TC,AZT,NVP 2003-12-01 2004-01-05      0      <NA>      NA      35
## 1002   3TC,AZT,NVP 2002-09-26 2004-03-29      0      <NA>      NA      365
## 1003   3TC,DDI,EFV 2007-01-31 2007-04-16      0      <NA>      NA      75
## 1004   3TC,D4T,NVP 2003-12-03 2007-10-11      0      <NA>      NA      365
##      lost init.reg_list 3TC  AZT  EFV  NVP  D4T  ABC  DDI  IDV  LPV  RTV
## 1      0 3TC, AZT, EFV 3TC  AZT  EFV <NA> <NA> <NA> <NA> <NA> <NA>
## 2      0 3TC, AZT, EFV 3TC  AZT  EFV <NA> <NA> <NA> <NA> <NA> <NA>
## 3      0 3TC, AZT, EFV 3TC  AZT  EFV <NA> <NA> <NA> <NA> <NA> <NA>
## 4      0 3TC, AZT, NVP 3TC  AZT <NA> NVP <NA> <NA> <NA> <NA> <NA>
## 5      0 3TC, D4T, EFV 3TC <NA> EFV <NA> D4T <NA> <NA> <NA> <NA>
## 1000   0 3TC, D4T, NVP 3TC <NA> <NA> NVP D4T <NA> <NA> <NA> <NA>
## 1001   1 3TC, AZT, NVP 3TC  AZT <NA> NVP <NA> <NA> <NA> <NA> <NA>
## 1002   0 3TC, AZT, NVP 3TC  AZT <NA> NVP <NA> <NA> <NA> <NA> <NA>
## 1003   1 3TC, DDI, EFV 3TC <NA> EFV <NA> <NA> <NA> DDI <NA> <NA> <NA>
## 1004   0 3TC, D4T, NVP 3TC <NA> <NA> NVP D4T <NA> <NA> <NA> <NA> <NA>
##      SQV  FTC  TDF  DDC  NFV  T20  ATV  FPV
## 1      <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 2      <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 3      <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 4      <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 5      <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 1000   <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 1001   <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 1002   <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 1003   <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
## 1004   <NA> <NA> <NA> <NA> <NA> <NA> <NA> <NA>
```

## Question 2

### 14 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {
  if(exists(".Random.seed", envir = .GlobalEnv)) {
    save.seed <- get(".Random.seed", envir= .GlobalEnv)
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))
  } else {
    on.exit(rm(".Random.seed", envir = .GlobalEnv))
  }
  set.seed(n)
  subj <- ceiling(n / 10)
  id <- sample(subj, n, replace=TRUE)
  times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))
  dt <- as.POSIXct(sample(times, n), origin='2000-01-01')
  mu <- runif(subj, 4, 10)
  a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)
  data.frame(id, dt, a1c)
```

```
}
x <- genData(500)
```

Perform the following manipulations: (2 points each)

1. Order the data set by id and dt.

```
x <- x[order(x$id,x$dt),]
```

2. For each id, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the a1c value set to missing. A two year gap would require two new rows, and so forth.

```
#Write a function that finds gaps
check.dates <- function(identity,date){
  insert.at <- vector()
  rows.fin <- vector()
  for (i in unique(identity)){
    rows <- which(identity==i)[1:length(which(identity==i))-1]
    for (j in rows){
      rows.fin <- c(rows.fin, j)
      if(unclass(difftime(date[j+1], date[j], "days"))[1] > 366){
        insert.at <- c(insert.at,j+1)
      }
    }
  }
  return(insert.at)
}

#Write a function that fills gaps
add.gap <- function(df,insertion){
  df <- rbind(df[1:(insertion-1),],data.frame(id=df$id[insertion-1],
                                              dt=df$dt[insertion-1]+years(1),a1c=NA),
             df[insertion:nrow(df),])
  return(df)
}

p <- x
insert.at <- check.dates(p$id,p$dt)
lines <- insert.at+seq(from=0,by=1,length.out=length(insert.at))
for (i in 1:length(lines)){
  p <- add.gap(p,lines[i])
}

#Check again to fix 2-year gaps
insert.at <- check.dates(p$id,p$dt)
lines <- insert.at+seq(from=0,by=1,length.out=length(insert.at))
for (i in 1:length(lines)){
  p <- add.gap(p,lines[i])
}

#Check for any 3-year gaps
(insert.at <- check.dates(p$id,p$dt))
```

```
## logical(0)
```

```
x <- p
```

3. Create a new column `visit`. For each `id`, add the visit number. This should be 1 to `n` where `n` is the number of observations for an individual. This should include the observations created with missing `a1c` values.

```
for (i in 1:length(unique(x$id))){  
  visits <- seq(1:table(x$id)[[i]])  
  x$visit[x$id==i] <- visits  
}
```

4. For each `id`, replace missing values with the mean `a1c` value for that individual.

```
for (i in 1:length(unique(x$id))){  
  rows <- which(x$id==i)  
  meana1c <- mean(x$a1c[rows[1]:tail(rows,n=1)],na.rm = TRUE)  
  for (j in rows){  
    if(is.na(x$a1c[j])){  
      x$a1c[j] <- meana1c  
    }  
  }  
}
```

5. Print mean `a1c` for each `id`.

```
for (i in 1:length(unique(x$id))){  
  rows <- which(x$id==i)  
  meana1c <- mean(x$a1c[rows[1]:tail(rows,n=1)])  
  print(c(as.integer(i),meana1c))  
}
```

```
## [1] 1.000000 4.063372  
## [1] 2.000000 7.544643  
## [1] 3.000000 6.75764  
## [1] 4.000000 3.892127  
## [1] 5.000000 9.512311  
## [1] 6.000000 7.555965  
## [1] 7.000000 9.161686  
## [1] 8.000000 7.189064  
## [1] 9.000000 9.283873  
## [1] 10.000000 7.975217  
## [1] 11.000000 6.917562  
## [1] 12.000000 7.034021  
## [1] 13.000000 9.145282  
## [1] 14.000000 6.623756  
## [1] 15.000000 8.012406  
## [1] 16.000000 4.222158  
## [1] 17.000000 3.996034  
## [1] 18.000000 9.164873
```

```
## [1] 19.00000 5.50721
## [1] 20.00000 3.726675
## [1] 21.00000 8.140939
## [1] 22.00000 5.637501
## [1] 23.00000 7.366889
## [1] 24.00000 7.439316
## [1] 25.00000 6.877135
## [1] 26.00000 6.556759
## [1] 27.00000 4.926457
## [1] 28.00000 7.433917
## [1] 29.00000 4.508086
## [1] 30.00000 6.045577
## [1] 31.00000 7.116586
## [1] 32.00000 6.568791
## [1] 33.00000 6.494069
## [1] 34.00000 6.768615
## [1] 35.0000 8.4767
## [1] 36.00000 9.60441
## [1] 37.00000 9.606253
## [1] 38.00000 5.355979
## [1] 39.00000 6.917013
## [1] 40.00000 9.530136
## [1] 41.00000 9.802424
## [1] 42.00000 3.89177
## [1] 43.00000 6.095849
## [1] 44.00000 9.09167
## [1] 45.00000 6.737204
## [1] 46.00000 9.621763
## [1] 47.00000 9.231489
## [1] 48.0000 6.4046
## [1] 49.00000 6.096076
## [1] 50.00000 8.962319
```

6. Print total number of visits for each id.

```
table(x$id)
```

```
##
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25
## 11 20 14 12 14 10  9 12 11 12 10 10  8 12  8  9 12 10 10  9 10  8  8 15 12
## 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50
## 14 11 14 10  7 11  5  8 12 11  9 17 15  8  7 17 14 11 11 14  9 12 11 12 10
```

7. Print the observations for id = 15.

```
x[which(x$id==15),]
```

```
##      id      dt      a1c visit
## 11   15 2000-04-30 00:34:50 7.527105    1
## 406  15 2001-01-17 21:11:02 5.898371    2
## 306  15 2001-04-25 06:23:05 8.566593    3
## 1117 15 2002-04-25 06:23:05 8.012406    4
```



```
## 1154 15 2003-04-25 06:23:05 8.012406      5
## 484  15 2003-06-06 14:06:00 9.133769      6
## 1118 15 2004-06-06 14:06:00 8.012406      7
## 263  15 2004-08-20 17:47:11 8.936190      8
```

### Question 3

#### 10 points

Import the `addr.txt` file from the GitHub repository. This file contains a listing of names and addresses (thanks google). Parse each line to create a data.frame with the following columns: `lastname`, `firstname`, `streetno`, `streetname`, `city`, `state`, `zip`. Keep middle initials or abbreviated names in the `firstname` column. Print out the entire data.frame.

```
url3 <- "https://raw.githubusercontent.com/fonnesbeck/Bios6301/master/datasets/addr.txt"
addr <- read.delim(url3, header=FALSE, stringsAsFactors=FALSE)
parsed.data <- data.frame(lastname=rep(NA,nrow(addr)),firstname=rep(NA,nrow(addr)),
                          streetno=rep(NA,nrow(addr)),streetname=rep(NA,nrow(addr)),
                          city=rep(NA,nrow(addr)),state=rep(NA,nrow(addr)),zip=rep(NA,nrow(addr)))

#Write function to trim leading and trailing whitespace to use after splitting strings into sections
trim <- function (x) gsub("^\\s+|\\s+$", "", x)

#Loop through addr file to parse into fields
for (i in 1:nrow(addr)){
  #Identify whitespace of 2 spaces or more
  cutpoints <- c(1,unlist(gregexpr(" {2,}", addr[i,])),nchar(addr[i,]))
  fields <- vector()
  #Loop through each line and cut into parts, then trim trailing and leading whitespace
  for (j in 1:(length(cutpoints)-1)){
    fields[j] <- substring(addr[i,],cutpoints[j],cutpoints[j+1])
    fields[j] <- trim(fields[j])
  }
  #Assign parts that don't need more splitting to the appropriate columns
  parsed.data[i,"lastname"] <- fields[1]
  parsed.data[i,"firstname"] <- fields[2]
  parsed.data[i,"city"] <- fields[4]
  parsed.data[i,"state"] <- fields[5]
  parsed.data[i,"zip"] <- fields[6]
  #Split street address into number and street name
  name.cut <- unlist(gregexpr("[[:alpha:]]", fields[3]))
  number <- substring(fields[3],1,name.cut[1]-1)
  parsed.data[i,"streetno"] <- trim(number)
  street <- substring(fields[3],name.cut[1],nchar(fields[3]))
  parsed.data[i,"streetname"] <- trim(street)
}
print(parsed.data)
```

```
##      lastname  firstname streetno      streetname      city state
## 1      Bania  Thomas M.      725  Commonwealth Ave.    Boston   MA
## 2    Barnaby    David      373      W. Geneva St.    Wms. Bay  WI
## 3    Bausch     Judy      373      W. Geneva St.    Wms. Bay  WI
## 4    Bolatto   Alberto      725  Commonwealth Ave.    Boston   MA
```

## 5	Carlstrom	John	933	E. 56th St.	Chicago	IL
## 6	Chamberlin	Richard A.	111	Nowelo St.	Hilo	HI
## 7	Chuss	Dave	2145	Sheridan Rd	Evanston	IL
## 8	Davis	E. J.	933	E. 56th St.	Chicago	IL
## 9	Depoy	Darren	174	W. 18th Ave.	Columbus	OH
## 10	Griffin	Greg	5000	Forbes Ave.	Pittsburgh	PA
## 11	Halvorsen	Nils	933	E. 56th St.	Chicago	IL
## 12	Harper	Al	373	W. Geneva St.	Wms. Bay	WI
## 13	Huang	Maohai	725	W. Commonwealth Ave.	Boston	MA
## 14	Ingalls	James G.	725	W. Commonwealth Ave.	Boston	MA
## 15	Jackson	James M.	725	W. Commonwealth Ave.	Boston	MA
## 16	Knudsen	Scott	373	W. Geneva St.	Wms. Bay	WI
## 17	Kovac	John	5640	S. Ellis Ave.	Chicago	IL
## 18	Landsberg	Randy	5640	S. Ellis Ave.	Chicago	IL
## 19	Lo	Kwok-Yung	1002	W. Green St.	Urbana	IL
## 20	Loewenstein	Robert F.	373	W. Geneva St.	Wms. Bay	WI
## 21	Lynch	John	4201	Wilson Blvd	Arlington	VA
## 22	Martini	Paul	174	W. 18th Ave.	Columbus	OH
## 23	Meyer	Stephan	933	E. 56th St.	Chicago	IL
## 24	Mrozek	Fred	373	W. Geneva St.	Wms. Bay	WI
## 25	Newcomb	Matt	5000	Forbes Ave.	Pittsburgh	PA
## 26	Novak	Giles	2145	Sheridan Rd	Evanston	IL
## 27	Odalen	Nancy	373	W. Geneva St.	Wms. Bay	WI
## 28	Pernic	Dave	373	W. Geneva St.	Wms. Bay	WI
## 29	Pernic	Bob	373	W. Geneva St.	Wms. Bay	WI
## 30	Peterson	Jeffrey	5000	Forbes Ave.	Pittsburgh	PA
## 31	Pryke	Clem	933	E. 56th St.	Chicago	IL
## 32	Rebull	Luisa	5640	S. Ellis Ave.	Chicago	IL
## 33	Renbarger	Thomas	2145	Sheridan Rd	Evanston	IL
## 34	Rottman	Joe	8730	W. Mountain View Ln	Littleton	CO
## 35	Schartman	Ethan	933	E. 56th St.	Chicago	IL
## 36	Spotz	Bob	373	W. Geneva St.	Wms. Bay	WI
## 37	Thoma	Mark	373	W. Geneva St.	Wms. Bay	WI
## 38	Walker	Chris	933	N. Cherry St.	Tucson	AZ
## 39	Wehrer	Cheryl	5000	Forbes Ave.	Pittsburgh	PA
## 40	Wirth	Jesse	373	W. Geneva St.	Wms. Bay	WI
## 41	Wright	Greg	791	Holmdel-Keyport Rd.	Holmdel	NY
## 42	Zingale	Michael	5640	S. Ellis Ave.	Chicago	IL
##	zip					
## 1	02215					
## 2	53191					
## 3	53191					
## 4	02215					
## 5	60637					
## 6	96720					
## 7	60208-3112					
## 8	60637					
## 9	43210					
## 10	15213					
## 11	60637					
## 12	53191					
## 13	02215					
## 14	02215					
## 15	02215					

```
## 16      53191
## 17      60637
## 18      60637
## 19      61801
## 20      53191
## 21      22230
## 22      43210
## 23      60637
## 24      53191
## 25      15213
## 26 60208-3112
## 27      53191
## 28      53191
## 29      53191
## 30      15213
## 31      60637
## 32      60637
## 33 60208-3112
## 34      80125
## 35      60637
## 36      53191
## 37      53191
## 38      85721
## 39      15213
## 40      53191
## 41 07733-1988
## 42      60637
```

#### Question 4

##### 2 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula.

```
url <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart_df <- read.csv(url)[,c('death','weight','hemoglobin','cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin   -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(response ~ ., data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is "catching" the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(expr, envir, enclos): object 'death' not found>
```

What do you think is going on? Consider using `debug` to trace the problem.

The function is failing because it cannot find the object 'death' in the function call. This occurs in the third line when we find the coefficients of the summary of the glm of `death ~ .` using `haart_df`. This is happening because the object "death" is defined inside of the data set `haart_df`, but we are trying to call the variable from outside of the data set itself. However, trying to call `haart_df$death` results in the model not properly using `death` as the outcome variable.

```
myfun(haart_df, haart_df$death)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
##           Estimate Std. Error      z value Pr(>|z|)
## (Intercept) -2.656607e+01 115935.1724 -2.291459e-04 0.9998172
## death        5.313214e+01  69028.4183  7.697140e-04 0.9993859
## weight      -4.499694e-15   1939.0571 -2.320558e-18 1.0000000
## hemoglobin   5.124642e-14   9774.8190  5.242697e-18 1.0000000
## cd4baseline  1.830771e-16    184.0846  9.945271e-19 1.0000000
```

## 5 bonus points

Create a working function.

*#We will use `deparse(substitute(x))` to convert the inputs into strings, create the formula manually, and*

```
myfun.2 <- function(dat, response) {
  c <- deparse(substitute(response))
  d <- deparse(substitute(dat))
  e <- paste(d, c, sep="$")
  f <- paste(e, " ~ .", sep="")
  print(f)
  print(coef(summary(glm(f, data=dat, family=binomial(logit)))))
}
```

```
myfun.2(haart_df, death)
```

```
## [1] "haart_df$death ~ ."
##           Estimate Std. Error      z value      Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```